

# **How can we improve Apprenticeship learning as DNN?**

Jeong Gwan Lee

KAIST

(Korea Advanced Institute of Science and Technology)

# Before I start,

---

In MDP, meaning of Reward function?

The reward function is "the most succinct, robust, and transferable definition of a task." [1]

[1] Abbeel, Pieter. 2008. Apprenticeship Learning and Reinforcement Learning with Application to Robotic Control. Ph.D. thesis, Stanford University, Stanford, CA, USA.

# Apprentice Learning,

---

is a learning approach, which assumes that **the reward function is unknown**, but instead, *expert knowledge is available*.

This *expert knowledge* is in the form of sequences of states and actions, trajectories, where the goal state is achieved.

The main idea is to find a optimal policy close to the *expert policy*  $\pi^E$ , to use these trajectories.

# Inverse Reinforcement Learning(IRL)

Reward function is parameterized as **linear combination** of features,

$$R(s) = w^T \phi(s) \quad s.t. \quad \phi : S \rightarrow [0, 1]^k,$$

w is the weight vector to be learned,  
 $\phi$  is the features vector(basis function),  
k is the number of features.

Feature expectations, denoted as  $\mu$ ,

$$\mu(\pi) = E\left[\sum_{t=0}^{\infty} \gamma^t \phi(s_t) | \pi\right] \in \mathbb{R}^k.$$

Estimated feature expectations for expert policy,

$$\hat{\mu}_E = \frac{1}{m} \sum_{i=0}^m \sum_{t=0}^{\infty} \gamma^t \phi(s_t^{(i)}),$$

Learning process would be repeated until,  $\|\mu(\pi) - \hat{\mu}_E\| < \epsilon$

# Batch, Off-Policy and Model-Free Apprenticeship Learning

# IRL set-up

The true reward function belongs to some hypothesis space

$$\mathcal{H}_\phi = \{\theta^T \phi(s), \theta \in \mathbb{R}^p\}, |\phi_i(s)| \leq 1, \forall s \in S, 1 \leq i \leq p.$$

$$R^*(s) = (\theta^*)^T \phi(s)$$

parameters, weights.

features; state representation; input.



$$R(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)$$

$f_i$  : DNN layer with activation function.

*Reward estimator*

$$\phi(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)$$



$f_i$  : DNN layer with activation function.

*Feature expect. estimator*

- 1) depending on input (videos or game display -> CNN, factory cases -> ?, gathering sensor input AMAP(temperature?)), 2) embedding space,
- 3) why R is related to s regardless of policy

# IRL set-up

For any reward function belonging to  $\mathcal{H}_\phi = \{\theta^T \phi(s), \theta \in \mathbb{R}^p\}$ ,  
Value function  $V(s)$  can be expressed,

$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t \theta^T \phi(s_t) \mid s_0 = s, \pi\right] = \theta^T E\left[\sum_{t=0}^{\infty} \gamma^t \phi(s_t) \mid s_0 = s, \pi\right]$$


1. (reward esti.) DNN, Non-linearity, so, might not make feature expectation?
2. (feature extractor)  $\phi(s_t)$  might be “the feature output just before softmax” in classification?  $\rightarrow$  Depending on the input.

*Feature expectation* is,

$$\mu^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t \phi(s_t) \mid s_0 = s, \pi\right]$$

$$|V^\pi(s_0) - V^{\tilde{\pi}}(s_0)| = |\theta^T (\mu^\pi(s_0) - \mu^{\tilde{\pi}}(s_0))| \leq \|\mu^\pi(s_0) - \mu^{\tilde{\pi}}(s_0)\|_2$$

# IRL Algorithm

1. Starts with some initial policy  $\pi^{(0)}$  and compute  $\mu^{\pi^{(0)}}(s_0)$ . Set  $j = 1$ ;
2. Compute  $t^{(j)} = \max_{\theta: \|\theta\|_2 \leq 1} \min_{k \in \{0, j-1\}} \theta^T (\mu^{\pi^E}(s_0) - \mu^{\pi^{(k)}}(s_0))$  and let  $\theta^{(j)}$  be the value attaining this maximum. At this step, one searches for the reward function which maximizes the distance between the value of the expert at  $s_0$  and the value of *any* policy computed so far (still at  $s_0$ ). This optimization problem can be solved using a quadratic programming approach or a projection algorithm [1];
3. if  $t^{(j)} \leq \epsilon$ , terminate. The algorithm outputs a set of policies  $\{\pi^{(0)}, \dots, \pi^{(j-1)}\}$  among which the user chooses manually or automatically the closest to the expert (see [1] for details on how to choose this policy). Notice that the last policy is not necessarily the best (as illustrated in Section 4);
4. solve the MDP with the reward function  $R^{(j)}(s) = (\theta^{(j)})^T \phi(s)$  and denote  $\pi^{(j)}$  the associated optimal policy. Compute  $\mu^{\pi^{(j)}}(s_0)$ ;
5. set  $j \leftarrow j + 1$  and go back to step 2.

LSPI



# LSTD- $\mu$

Compute  $\mu^{\pi^{(j)}}(s_0)$

LSTD- $\mu$  : to estimate feature expectation of intermediate policies

$$\mu_i^{\pi}(s) = E\left[\sum_{t=0}^{\infty} \gamma^t \phi_i(s_t) | s_0 = s, \pi\right].$$

seems like Value function

$$\mathcal{H}_{\psi} = \{\hat{V}_{\xi}(s) = \sum_{i=1}^q \xi_i \psi_i(s) = \xi^T \psi(s), \xi \in \mathbb{R}^q\}$$

LSTD estimate

$$\xi_i^* = \left( \sum_{t=1}^n \psi(s_t) (\psi(s_t) - \gamma \psi(s'_t))^T \right)^{-1} \sum_{t=1}^n \psi(s_t) \phi_i(s_t)$$

$$(\hat{\mu}^{\pi}(s_0))^T = \psi(s_0)^T (\Psi^T \Delta \Psi)^{-1} \Psi^T \Phi$$

# LSTD- $\mu$

Compute  $\mu^{\pi^{(j)}}(s_0)$

LSTD- $\mu$  : to estimate feature expectation of intermediate policies

$$\mu_i^{\pi}(s) = E\left[\sum_{t=0}^{\infty} \gamma^t \phi_i(s_t) | s_0 = s, \pi\right].$$

psi  $\approx$  Reward function

seems like Value function



$$\mathcal{H}_{\psi} = \{\hat{V}_{\xi}(s) = \sum_{i=1}^q \xi_i \psi_i(s) = \xi^T \psi(s), \xi \in \mathbb{R}^q\}$$

$$V(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)?$$

$$\psi(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)?$$

$$\xi_i^* = \left( \sum_{t=1}^n \psi(s_t) (\psi(s_t) - \gamma \psi(s'_t))^T \right)^{-1} \sum_{t=1}^n \psi(s_t) \phi_i(s_t) \quad \text{LSTD estimate}$$


$$V(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)?$$

Efficient way to estimate the feature expectation of the expert in  $s_0$

# LSTD- $\mu$ as off-policy manner

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left( R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right) \quad \text{Q-learning as off-policy algorithm}$$

State-Action Feature expectation  $\mu^\pi(s, a) = E\left[\sum_{t=0}^{\infty} \gamma^t \phi(s_t) \mid s_0 = s, a_0 = a, \pi\right]$



regards as state-action Q-function

LSTD-Q : learn a policy  $\pi$  by estimating a linear approximation  $\hat{Q}^\pi = \Phi \omega^\pi$   
 $A \omega^\pi = b$ , where  $A = \Phi^T (\Phi - \gamma \mathbf{P} \Pi_\pi \Phi)$  and  $b = \Phi^T \mathcal{R}$ .

Additional degree of freedom ( $a_0 = a$ ) allows off-policy learning. (LSTD-Q)  
LSTD-Q (Q-function)  $\rightarrow$  LSTD- $\mu$  (state-action feature expectation)

# Deep LSTD- $\mu$

## 1) Deep feature expectation for Reward hypothesis space

$$R^*(s) = (\theta^*)^T \phi(s) \quad \phi(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)$$

$f_i$  : DNN layer with activation function.

$$\mu^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t \phi(s_t) \mid s_0 = s, \pi\right]$$

## 2) Deep feature extraction for $\mu$ hypothesis space

$$\mathcal{H}_\psi = \{\hat{V}_\xi(s) = \sum_{i=1}^q \xi_i \psi_i(s) = \xi^T \psi(s), \xi \in \mathbb{R}^q\} \quad \psi(s) = f_N(\dots (f_2(f_1(x, \theta_1), \theta_2), \dots), \theta_N)$$

$$\xi_i^* = \left( \sum_{t=1}^n \psi(s_t) (\psi(s_t) - \gamma \psi(s'_t))^T \right)^{-1} \sum_{t=1}^n \psi(s_t) \phi_i(s_t)$$

# My Conclusion

---

What is important in IRL is,

1) Expert trajectories, 2) semantic inputs for Reward function

How to get Reward function in IRL are,

1) approximating R well or 2) narrowing the boundary of R

Input(states, features) and Reward function are closely connected in IRL.



DNN might help to infer the reward function.

Lots of complex, meaningful inputs are needed.

# My Conclusion

---

Limitation of IRL is,

1) need expert trajectories, so it can't self-study.

Human can do → machine can do!

Human can't do → machine can't do 😞